

Gradient Descent = not efficient !

Sometimes learning rate too small  $\Rightarrow$  slow

Sometimes learning rate too large  $\Rightarrow$  diverges and doesn't reach minimum

There is an algorithm called the ADAM algorithm



Adaptive Movement Estimation



Adjusts 'α' (the learning rate) accordingly

how?



doesn't use 1 α

Uses different α for each parameter ( $\vec{w}$  and  $b$ )

For eg. if our algorithm has 10  $w$ 's.

$$w_1 = w_1 - \alpha_1 \frac{\partial J(\vec{w}, b)}{\partial w_1}$$


.

.


$$w_{10} = w_{10} - \alpha_{10} \frac{\partial J(\vec{w}, b)}{\partial w_{10}}$$

$$b = b - \alpha_{11} \frac{\partial J(\vec{w}, b)}{\partial b}$$

## Intuition of ADAM algorithm



If  $w_j$  (or  $b_j$ ) moves in the same direction then increase  $\alpha_j$ .



If  $w_j$  (or  $b_j$ ) keeps oscillating, then reduce  $\alpha_j$ .

## IN CODE

### model

```
model = Sequential([
    Dense(units=25, activation="relu"),
    Dense(units=15, activation="relu"),
    Dense(units=1, activation="linear")
])
```

### compile

new line

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),  
               $\alpha = 10^{-3} = 0.001$  (just needs a default value),  
              loss=tf.keras.losses.SparseCategoricalCrossEntropy(from_logits=True))
```

fit

```
model.fit(x, y, epochs = 100)
```