

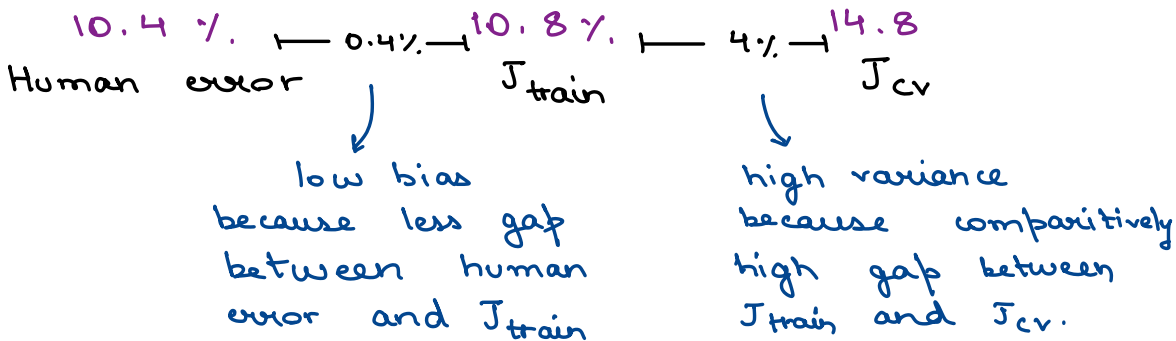
The need to have a baseline for bias and variance

Example, take speech recognition, it has J_{train} error: 10.8% and J_{cv} : 14.8%.

We might say it has high bias and high variance (i.e. underfit), but we have to be realistic with our data.

Some audio on the internet is intelligible, even humans make 10.4% error while converting speech to text.

If even humans make that much error then it's safe to say the algorithm performs very good.



Basically,

1. Human error \leftarrow less gap $\rightarrow J_{\text{train}}$ \leftarrow high gap $\rightarrow J_{\text{cv}}$

eg. $10.4\% \longleftrightarrow J_{\text{train}}: 10.8\% \longleftrightarrow J_{\text{cv}}: 14.8\%$
 \Rightarrow High variance

2. Human error \leftarrow high gap $\rightarrow J_{\text{train}}$ \leftarrow less gap $\rightarrow J_{\text{cv}}$

eg. $10.4\% \longleftrightarrow J_{\text{train}}: 14.8\% \longleftrightarrow J_{\text{cv}}: 15\%$
 \Rightarrow High bias

3. Human error \leftarrow high gap $\rightarrow J_{\text{train}}$ \leftarrow high gap $\rightarrow J_{\text{cv}}$

eg. $10.4\% \longleftrightarrow J_{\text{train}}: 14.8\% \longleftrightarrow J_{\text{cv}}: 18\%$
 \Rightarrow High variance and high bias

Establishing a baseline level of performance

How to obtain the level of error we can possibly hope to get to?

\rightarrow Comparing to human level performance

\rightarrow Comparing other algorithms' performance

\rightarrow Take a guess based on prior experience